

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
df = pd.read_csv('Advertising.csv')
```

```
print(df.head(5))
```

```
↗
```

	Unnamed: 0	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9

```
print(df.info())
```

```
↗
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0    200 non-null    int64
1   TV            200 non-null    float64
2   Radio         200 non-null    float64
3   Newspaper     200 non-null    float64
4   Sales         200 non-null    float64
dtypes: float64(4), int64(1)
memory usage: 7.9 KB
None
```

```
print(df.describe())
```

```
↗
```

	Unnamed: 0	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000	200.000000
mean	100.500000	147.042500	23.264000	30.554000	14.022500
std	57.879185	85.854236	14.846809	21.778621	5.217457
min	1.000000	0.700000	0.000000	0.300000	1.600000
25%	50.750000	74.375000	9.975000	12.750000	10.375000
50%	100.500000	149.750000	22.900000	25.750000	12.900000
75%	150.250000	218.825000	36.525000	45.100000	17.400000
max	200.000000	296.400000	49.600000	114.000000	27.000000

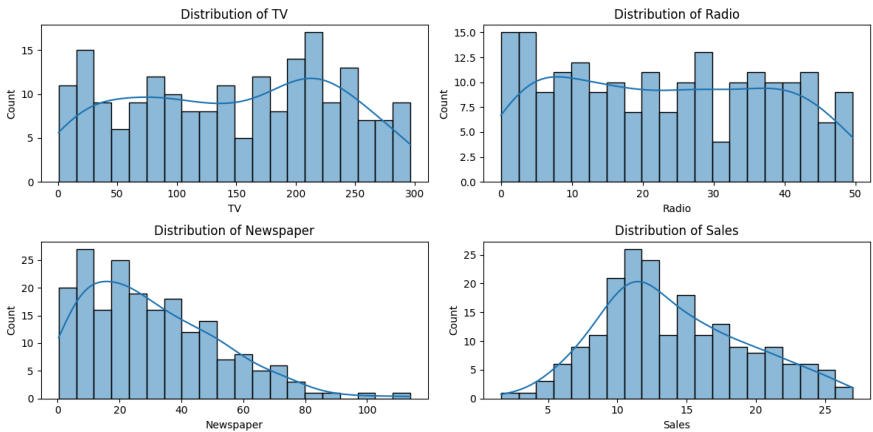
```
print(df.isnull().sum())
```

```
↗
```


```
Unnamed: 0    0
TV            0
Radio         0
```

```
Newspaper    0
Sales        0
dtype: int64
```

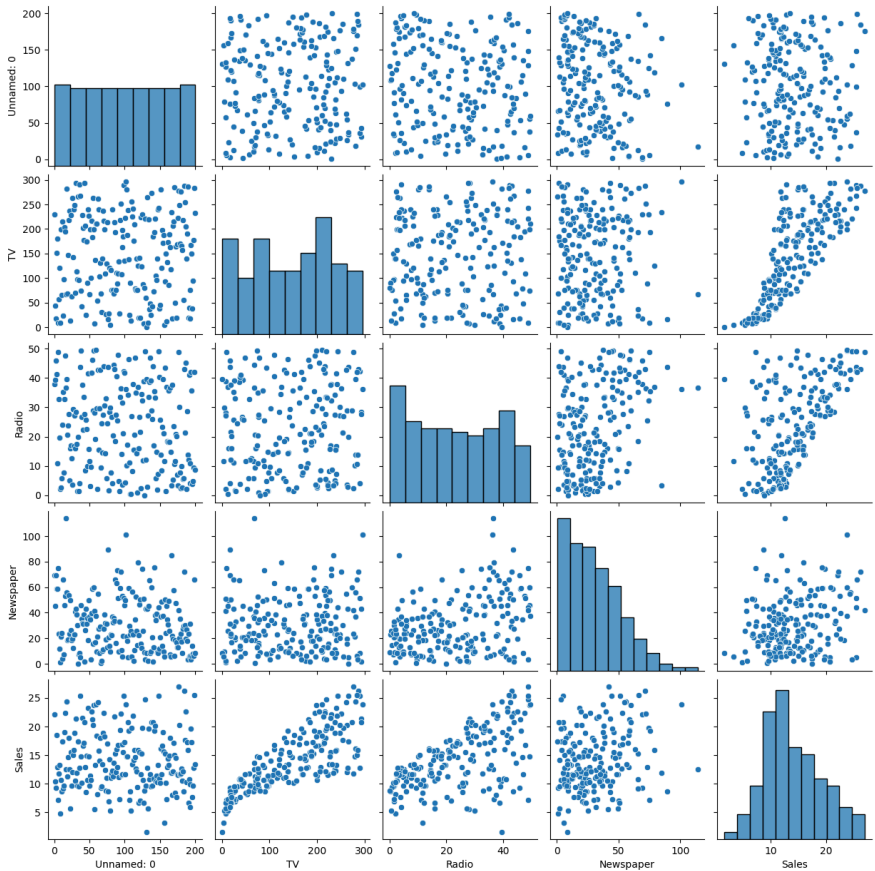
```
# Visualize the distribution of features
plt.figure(figsize=(12, 6))
for i, column in enumerate(['TV', 'Radio', 'Newspaper', 'Sales'], 1):
    plt.subplot(2, 2, i)
    sns.histplot(df[column], kde=True, bins=20)
    plt.title(f'Distribution of {column}')
plt.tight_layout()
plt.show()
```



```
# Correlation matrix
correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



```
# Pairplot to visualize relationships between variables
sns.pairplot(df)
plt.show()
```



```
df.drop(columns=['Unnamed: 0'], inplace=True, errors='ignore')
```

```
# Define features (X) and target variable (y)
X = df[['TV', 'Radio', 'Newspaper']]
y = df['Sales']
```

```
# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=
```

```
# Check the shapes of the datasets
print("Training set shape:", X_train.shape, y_train.shape)
print("Testing set shape:", X_test.shape, y_test.shape)
```

```
↗ Training set shape: (160, 3) (160,)
   Testing set shape: (40, 3) (40,)
```

```
model = LinearRegression()
```

```
# Train the model on the training data
model.fit(X_train, y_train)
```

```
↗ LinearRegression ⓘ ?
```

```
# Print the coefficients and intercept
print("Coefficients:", model.coef_)
print("Intercept:", model.intercept_)
```

```
↗ Coefficients: [0.04472952 0.18919505 0.00276111]
   Intercept: 2.979067338122629
```


```
y_pred = model.predict(X_test)
```

```
# Calculate evaluation metrics
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)
```

```
# Print the evaluation metrics
print("Mean Squared Error (MSE):", mse)
print("Root Mean Squared Error (RMSE):", rmse)
print("R-squared (R2):", r2)
```

```
↗ Mean Squared Error (MSE): 3.1740973539761033
   Root Mean Squared Error (RMSE): 1.78159966153345
   R-squared (R2): 0.899438024100912
```

```
# Visualize actual vs predicted values
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red', line:
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Actual vs Predicted Sales')
plt.show()
```



```
# Create a DataFrame to display feature importance
feature_importance = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': model.coef_
}).sort_values(by='Coefficient', ascending=False)

print(feature_importance)
```



	Feature	Coefficient
1	Radio	0.189195
0	TV	0.044730
2	Newspaper	0.002761

```
import joblib
```